

# Enabling Fastai Multi-GPU/DDP Training in Jupyter

- **fastai**: to make building high-performing AI application easy
  - Open source, free lessons videos + Jupyter notebooks
- **Distributed Data Parallel (DDP)**:
  - *a multiprocessing-based parallelism to speed up training with multiple GPUs/nodes.*

"fastai + DDP" cmdline app is trivial [1], but ...

*"Distributed training **doesn't** work in a notebook..."*[2]

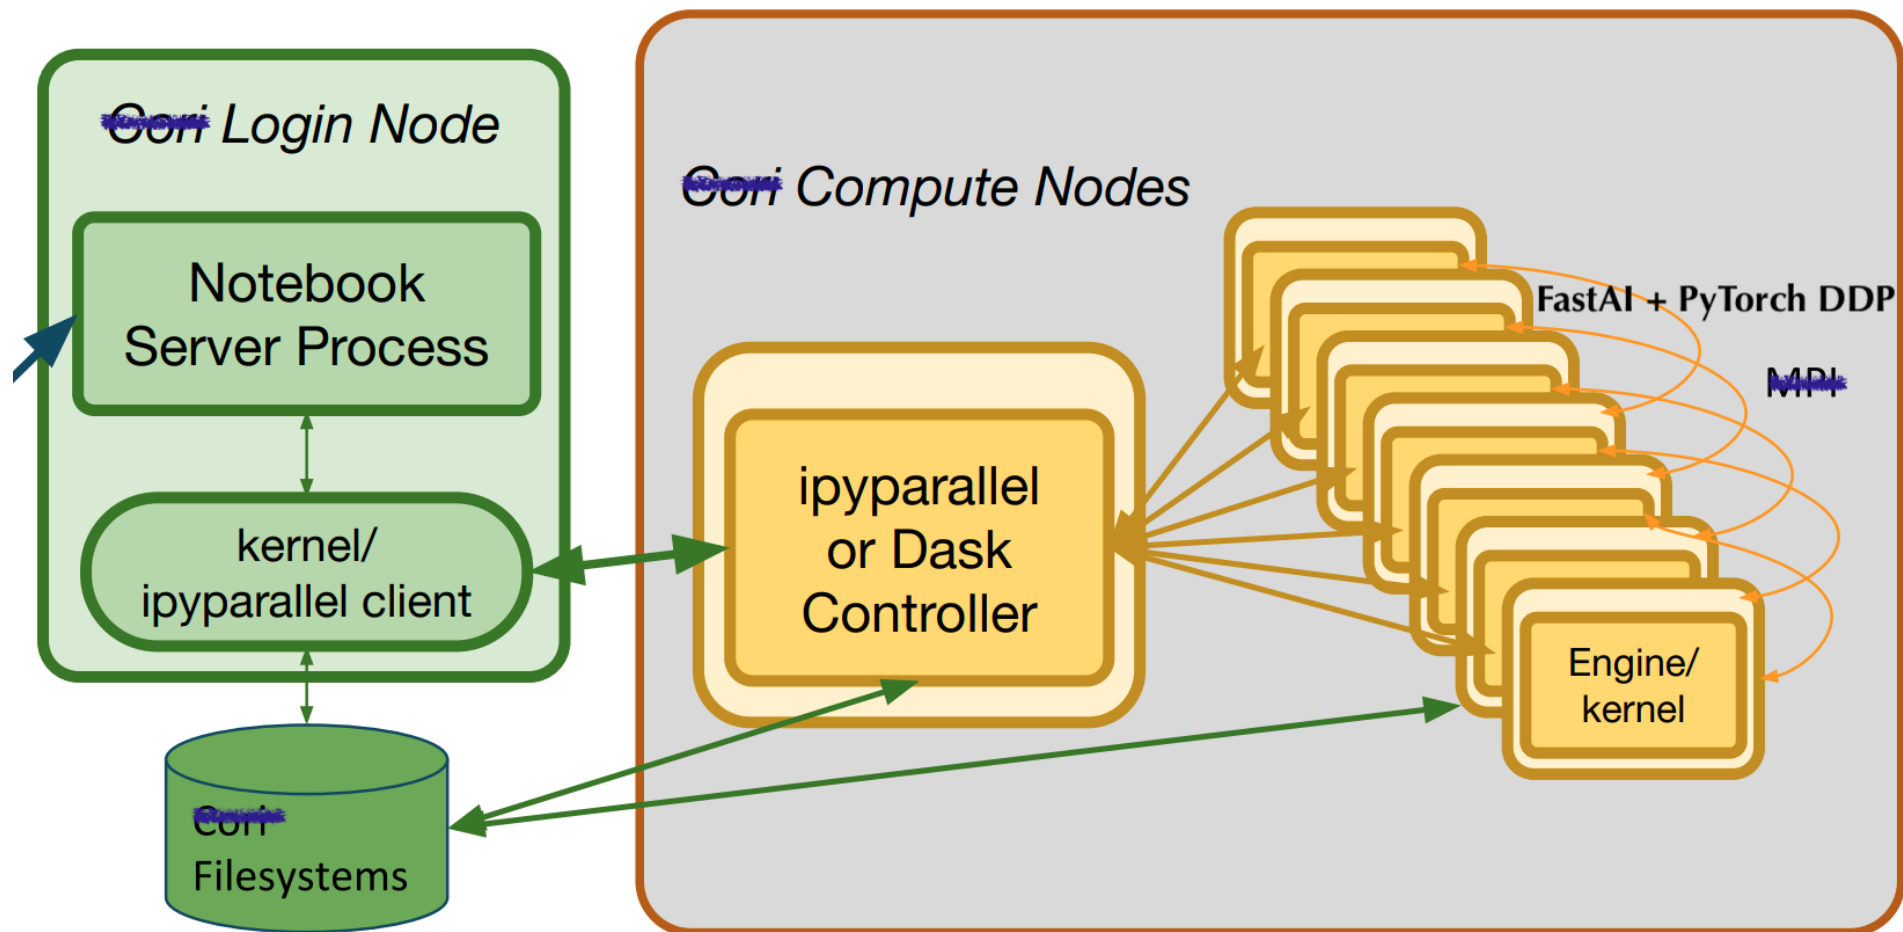
**Ddip** tries to bridge this gap.

[1] See [Reproducing DAWNBench winning results in a few lines of code](#)

[2] FastAI's tutorial on [How to launch a distributed training](#)

## ddp - Distributed data "interactive" parallel

An iPython extension to control PyTorch DDP from within Jupyter, uses `ipyparallel` underneath.



# `ddp` - Distributed data "interactive" parallel

---

## [Overview and Usage](#)

## [Speedup in Training](#)

### Limitations

- Works a single host only. Luckily `ipyparallel` does support cluster of nodes.
- Not all models gain: one model is flat, one has accuracy problem.
- *a wgan model manages to achieve linear slow down!!* 😞

# **Ddip** - Distributed data "interactive" parallel

---

## Lessons learned:

1. Python's dynamic nature empowers: patch classes dynamically, toss objects/functions across multiple processes.
2. Hooks/callback architecture brings openness and flexibility: Jupyter, `fastai`.
3. Multiprocess, multi-GPUs and Jupyter: both **Complexity and Opportunities**:
  - data movement, race conditions, proc & mem mgmt.
4. Design choices: semantics (`%`, `%%` vs library calls), deadlock solution, implicit convenience.

## Looking Forward

- `fastai v2` is out already.
- `nbdev` as productivity boost.
- Support cluster of nodes.
- Feedbacks and contribution via github are welcome!

Github Repo: <https://github.com/philtrade/Ddip/>